

SCRIBE: Consistent Distributed Environment Modeling for Communication-Limited Robotic Exploration

Shashank Swaminathan^{1,2} and Brian Williams²

Abstract—SCRIBE addresses distributed environment-model learning under limited communication while preserving accurate, mutually consistent estimates of the field and its uncertainty. It represents spatial models as fixed-dimensional linear states and introduces a fusion process that lets agents repeatedly assimilate any available peer’s model without full-team participation. Information Kalman filtering adds new evidence, Covariance Intersection reconciles divergent priors, and Metropolis–Hastings consensus propagates innovations using only local neighbor information. The resulting model update is lightweight, asynchronous, and consistent under intermittent communication. In 96 missions with 3–10 agents, final normalized error is 0.405–0.409, against 0.582–1.227 without communication. Across 432 planning trials, median fusion through ten agents remains below 6.2 ms and 24 kB per update. Independent Vulcan planners achieve pooled posterior variance within 3.2% of ideal centralized sharing, demonstrating effective exploration through communication loss.

Index Terms—autonomous exploration, distributed estimation, informative path planning, intermittent communication, multi-robot systems

I. INTRODUCTION

Autonomous exploration uses an uncertain environment model to decide which observations are most valuable [1], [2]. This is particularly important in the deep ocean, where vehicles may localize plumes, characterize cold seeps, or follow chemical gradients from sparse data [3]. Teams expand coverage, but acoustic networks are intermittent, bandwidth-limited, and range-dependent [1]. Robots must therefore learn and plan in isolation, then incorporate peer knowledge whenever contact resumes.

Distributing environment-model learning is non-trivial: each robot requires a locally useful model, yet team models must remain consistent in prediction and confidence. Intermittent connectivity causes robots to accumulate different evidence and their learned states to diverge. Existing methods either exchange observations and relearn from the collective history, whose cost grows with mission duration and sensor dimension [4], [5], or exchange and average compact model parameters, which can count correlated information more than once and overstate certainty [6]. The fundamental problem is not communication reduction alone, but guaranteeing an accurate, consistently calibrated shared model when only subsets of the team can participate in each update.

¹Computer Science and Artificial Intelligence Laboratory, Massachusetts Institute of Technology, Cambridge, MA, USA {shswami, williams}@mit.edu

²Applied Ocean Science & Engineering Department, Woods Hole Oceanographic Institution, Woods Hole, MA, USA shashank.swaminathan@whoi.edu




Fig. 1. Two SCRIBE models. Top: a simulated field, one robot’s final GSF estimate, and four robot paths. Bottom: held-out Ram Head ROMS snapshot 5021 and its EOF estimate after 300 samples ($b = 2$, RMSE 0.049).

Work on distributed modeling via GP updates [7], finite-dimensional regression [8], and the use of those shared models [9], [10], [11] demonstrates promising work in distributing environment models used for exploration, yet does not address the concern of maintaining model consistency, particularly in communication-limited conditions. SCRIBE closes this gap by contributing a model-learning and fusion mechanism that remains consistent under intermittent communication. A shared GSF or EOF basis first recasts the spatial model as a fixed-dimensional linear coefficient state. SCRIBE then interleaves an Information Kalman Filter, whose observation updates are additive, with distributed Covariance Intersection, which reconciles divergent priors without their unknown cross-correlations [6], [12]. Metropolis–Hastings consensus propagates the new information using only current neighbor degrees [13]. This hybrid construction [14] allows any available agents to repeatedly fuse their learned fields without loss of consistency, a global graph, or a synchronized team-wide update.

II. DISTRIBUTED ENVIRONMENT MODELING

A. Fixed-Dimensional Information State

When no collection of earlier field maps is available, SCRIBE represents a scalar field with n_ϕ Gaussian scalar-field (GSF) basis functions [8], [15],

$$\mathcal{F}(x, k) = \psi^T(x) \phi(k), \quad \psi_i(x) = \tau_i^{-1} e^{-\|x - \mu_i\|_2^2 / \sigma_i^2}. \quad (1)$$

The centers μ_i and widths σ_i place overlapping functions across the region; the coefficients $\phi_i(k)$ determine how




Fig. 2. Twelve four-robot trials of 36 rounds. Shading marks complete blackout; otherwise each robot communicates with at most one in-range peer. Curves give the mean and one standard deviation. The right panel rebuilds a model from all samples chosen by each team and therefore compares path quality rather than model fusion. The disagreement axis is truncated; lower is better.

strongly each function appears in the field at time k .

When a prior environmental record is available, a mean-centered, spatially weighted singular value decomposition instead yields empirical orthogonal function (EOF) modes [16]. If $\mathbf{E}_r$ contains the r retained modes at the stored feature locations, the field vector is

$$\mathbf{f}(k) = \bar{\mathbf{f}} + \mathbf{E}_r \phi(k), \quad (2)$$

where the prior mean $\bar{\mathbf{f}}$ and modes remain fixed while $\phi(k)$ describes the current field. GSFs suit a poorly known environment, while EOFs compress spatial structure present in earlier data.

Online estimation differs only in constructing the observation matrix. GSF rows are $\psi^T(x)$ at the sample locations. For EOFs, interpolation $\mathbf{L}_{\mathbf{X}_j}$ maps the stored grid to those locations, giving $\mathbf{H}_j = \mathbf{L}_{\mathbf{X}_j} \mathbf{E}_r$ and centered data $\mathbf{z}'_j = \mathbf{z}_j - \mathbf{L}_{\mathbf{X}_j} \bar{\mathbf{f}}$. GSF data require no mean subtraction. Both satisfy $\mathbf{z}'_j = \mathbf{H}_j \phi + \mathbf{v}_j$. Agent j stores $\mathbf{Y}_j = \mathbf{P}_j^{-1}$ and $\mathbf{y}_j = \mathbf{Y}_j \hat{\phi}_j$; a measurement adds

$$\delta \mathbf{I}_j = \mathbf{H}_j^T \mathbf{R}_j^{-1} \mathbf{H}_j, \quad \delta \mathbf{i}_j = \mathbf{H}_j^T \mathbf{R}_j^{-1} \mathbf{z}'_j. \quad (3)$$

SCRIBE exchanges an $O(n_\phi^2)$ state independent of mission duration, rather than an $O(kd_z)$ observation history. From $\hat{\phi}_j = \mathbf{Y}_j^{-1} \mathbf{y}_j$, GSFs predict $\psi^T(x) \hat{\phi}_j$; EOFs use $\mathbf{L}_\mathbf{X}(\bar{\mathbf{f}} + \mathbf{E}_r \hat{\phi}_j)$.

B. Fusion under an Unstable Network

Agents' priors may contain shared observations, so adding them can count the same evidence repeatedly. Let $\mathcal{N}_j^{\ell+} = \mathcal{N}_j^\ell \cup \{j\}$ and $\Delta_j^\ell = \{\omega \geq 0 : \sum_{m \in \mathcal{N}_j^{\ell+}} \omega_m = 1\}$. Distributed CI fuses these correlated priors without requiring their cross-correlations:

$$\omega^* = \arg \min_{\omega \in \Delta_j^\ell} J \left[\left(\sum_{m \in \mathcal{N}_j^{\ell+}} \omega_m \mathbf{Y}_m^\ell \right)^{-1} \right], \quad (4)$$

$$(\tilde{\mathbf{y}}_j^{\ell+1}, \tilde{\mathbf{Y}}_j^{\ell+1}) = \sum_{m \in \mathcal{N}_j^{\ell+}} \omega_m^* (\mathbf{y}_m^\ell, \mathbf{Y}_m^\ell). \quad (5)$$

Current measurement innovations are new evidence and are instead averaged by MHMC. For $\delta \mathbf{s}_j \in \{\delta \mathbf{i}_j, \delta \mathbf{I}_j\}$,

$$\gamma_{jm}^\ell = \frac{1}{1 + \max(|\mathcal{N}_j^\ell|, |\mathcal{N}_m^\ell|)}, \quad (6)$$

$$\delta \mathbf{s}_j^{\ell+1} = \delta \mathbf{s}_j^\ell + \sum_{m \in \mathcal{N}_j^\ell} \gamma_{jm}^\ell (\delta \mathbf{s}_m^\ell - \delta \mathbf{s}_j^\ell). \quad (7)$$

The weights require only current neighbor degrees. After consensus within a component of size n_c ,

$$\mathbf{y}_j^+ = \tilde{\mathbf{y}}_j + n_c \delta \tilde{\mathbf{i}}_j, \quad \mathbf{Y}_j^+ = \tilde{\mathbf{Y}}_j + n_c \delta \tilde{\mathbf{I}}_j. \quad (8)$$

CI therefore preserves consistent confidence while MHMC distributes new observations. Fusion begins whenever a connection appears, without a global graph or common schedule. *Computational scaling.* For basis size b , messages are $\Theta(b^2)$ and assimilating q_j samples is $O(q_j b^2)$. An MHMC round is $O(d_j b^2)$ locally and $O(E b^2)$ fleet-wide; CI is $O(J_{\text{CI}}(d_j b^2 + b^3))$. Hence bounded-degree fleets scale as $O(N b^2)$ per round, versus $O(N^2 b^2)$ for complete graphs. Fixed-graph convergence is $O(\log(e_0/\epsilon)/(1-\rho))$; for changing networks, it depends on the connection process.

III. EXPERIMENTS AND RESULTS

Consistency after reconnection. Four agents first follow matched paths, isolating fusion from planning. Across 30 trials with disjoint networks, a middle-third blackout, and reconnection, SCRIBE reaches RMSE 0.117 ± 0.011 , close to centralized sharing (0.101 ± 0.012) and CI (0.118 ± 0.012), but below IKF without CI (0.288 ± 0.020) and no communication (0.814 ± 0.018). Fig. 1 also shows that EOFs trained on prior Ram Head ROMS output [17] retain the dominant held-out structure; the two-mode estimate reaches RMSE 0.049 after 300 samples.

Fusion without scheduled communication. Agents then sense continuously for 8178 steps in adjacent regions and exchange models only at proximity encounters. Every team size in Fig. 3 converges near centralized NRMSE (0.404–0.409), versus 0.582–1.227 without communication. The six-agent trajectories converge from distinct initial errors; residual spatial smoothing is the two-mode representation floor, not an asynchronous-fusion failure.




Fig. 3. Asynchronous fusion during continuous sensing. Circles mark peer encounters in the selected six-agent mission; the bottom row compares its final model with the ROMS ground truth.

Operational scaling. Across 432 trials with 2–10 agents, Fig. 4 separates estimation quality from fusion cost. Posterior variance remains stable, while network mixing governs consensus rounds, time, and payload. The worst case, the intermittent moderate-degree graph at ten agents, requires about 26 rounds, 6.2 ms, and 24 kB per update. Model quality therefore persists as the fleet grows; connectivity determines the operational cost.

Effect on information-guided exploration. Four independent Vulcan planners [18] explore for 36 rounds with distance-limited networks and a twelve-round blackout in Fig. 2. SCRIBE has the lowest final distributed error and, like CI, contracts disagreement after reconnection; IKF remains high because it does not reconcile correlated priors. SCRIBE’s final RMSE approaches centralized sharing, improves on CI, and halves the no-communication error. Pooling each regime’s samples indicates the quality of the resultant paths under the models learned under the sharing regime. SCRIBE nearly matches centralized pooled RMSE and integrated variance.

IV. CONCLUSION

SCRIBE provides consistent, fixed-dimensional fusion through intermittent communication. Tests through ten agents approach centralized information collection, while the integrated Kalman-Filter and Covariance Intersection approach keeps local models both usable during blackouts and from diverging with other agent models after regaining contact. The modeling approach effectively supports distributed exploration and planning like seen in [9], [19], and leaves integration with these planners as the next step.

REFERENCES

- [1] J. Hwang, N. Bose, and S. Fan, “AUV adaptive sampling methods: A review,” *Applied Sciences*, vol. 9, no. 15, p. 3145, 2019.
- [2] R. Cui, Y. Li, and W. Yan, “Mutual information-based multi-AUV path planning for scalar field sampling using multidimensional RRT,” *IEEE Trans. Syst., Man, Cybern.: Syst.*, vol. 46, no. 7, pp. 993–1004, 2016.
- [3] P. Vrolijk *et al.*, “Using a ladder of seeps with computer decision processes to explore for and evaluate cold seeps on the Costa Rica active margin,” *Frontiers in Earth Science*, vol. 9, p. 601019, 2021.




Fig. 4. Uncertainty, computation, and communication scaling across team sizes, network densities, and connectivity conditions. Dashed curves denote 50% communication availability.

- [4] A. Viseras *et al.*, “Decentralized multi-agent exploration with online-learning of Gaussian processes,” in *Proc. IEEE Int. Conf. Robot. Autom. (ICRA)*, 2016, pp. 4222–4229.
- [5] R. Allamraju and G. Chowdhary, “Communication efficient decentralized Gaussian process fusion for multi-UAS path planning,” in *Proc. Amer. Control Conf. (ACC)*, 2017, pp. 4442–4447.
- [6] L. Chen, P. O. Arambel, and R. K. Mehra, “Estimation under unknown correlation: Covariance intersection revisited,” *IEEE Trans. Autom. Control*, vol. 47, no. 11, pp. 1879–1882, 2002.
- [7] D. Jang, J. Yoo, C. Y. Son, D. Kim, and H. J. Kim, “Multi-robot active sensing and environmental model learning with distributed Gaussian process,” *IEEE Robot. Autom. Lett.*, vol. 5, no. 4, pp. 5905–5912, 2020.
- [8] G. Pilonetto, L. Schenato, and D. Varagnolo, “Distributed multi-agent Gaussian regression via finite-dimensional approximations,” *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 41, no. 9, pp. 2098–2111, 2019.
- [9] D. Jang, J. Yoo, C. Y. Son, and H. J. Kim, “Fully distributed informative planning for environmental learning with multi-robot systems,” *arXiv preprint arXiv:2112.14433*, 2021.
- [10] B. Nguyen *et al.*, “Spatially temporally distributed informative path planning for multi-robot systems,” in *Proc. Amer. Control Conf. (ACC)*, 2025, pp. 3429–3434.
- [11] J. Kim, D. Jang, and H. J. Kim, “Distributed multi-agent target search and tracking with Gaussian process and reinforcement learning,” *Int. J. Control Autom. Syst.*, vol. 21, no. 9, pp. 3057–3067, 2023.
- [12] Y. Wang and X. R. Li, “Distributed estimation fusion with unavailable cross-correlation,” *IEEE Trans. Aerosp. Electron. Syst.*, vol. 48, no. 1, pp. 259–278, 2012.
- [13] L. Xiao, S. Boyd, and S. Lall, “A scheme for robust distributed sensor fusion based on average consensus,” in *Proc. 4th Int. Symp. Inf. Process. Sensor Netw. (IPSN)*, 2005, pp. 63–70.
- [14] A. Tamjidi, R. Oftadeh, M. N. G. Mohamed, D. Yu, S. Chakravorty, and D. Shell, “Unifying consensus and covariance intersection for efficient distributed state estimation over unreliable networks,” *IEEE Trans. Robot.*, vol. 37, no. 5, pp. 1525–1538, 2021.
- [15] M. Jadaliha and J. Choi, “Environmental monitoring using autonomous aquatic robots: Sampling algorithms and experiments,” *IEEE Trans. Control Syst. Technol.*, vol. 21, no. 3, pp. 899–905, 2013.
- [16] R. Young, S. Merrifield, M. Anderson, M. Mazloff, and E. Terrill, “A greedy depth-seeking behavior for energy-efficient transits by an autonomous underwater vehicle,” *IEEE J. Ocean. Eng.*, vol. 49, no. 4, pp. 1383–1396, 2024.
- [17] W. Zhang, Y. Jia, A. Apprill, and T. A. Mooney, “Fine-scale hydrodynamics around St. John, US Virgin Islands. Part I: Spatial and temporal heterogeneity in the coastal environment,” *Frontiers in Marine Science*, vol. 12, p. 1464627, 2025.
- [18] B. J. Ayton and B. C. Williams, “Vulcan: A Monte Carlo algorithm for large chance constrained MDPs with risk bounding functions,” *arXiv preprint arXiv:1809.01220*, 2018.
- [19] J. Olkin, V. Parimi, and B. C. Williams, “Multi-agent Vulcan: An information-driven multi-agent path finding approach,” in *Proc. IEEE/RSJ Int. Conf. Intell. Robots Syst. (IROS)*, 2024, pp. 10 253–10 259.